

Exhibit 3

Home (/en-au/) / Blog (/en-au/blog/) / Announcements

Cloud storage now more affordable: Announcing general availability of Azure Archive Storage

(facebook.com/share.php?u=https%3A%2F%2Fazure.microsoft.com%2Fblog%2Fcloud-storage-now-more-affordable-general-availability-of-azure-archive-storage%2F)

(com/share?url=https%3A%2F%2Fazure.microsoft.com%2Fblog%2Fcloud-storage-now-more-affordable-general-availability-of-azure-archive-

(Cloud+storage+now+more+affordable%3A+Announcing+general+availability+of+Azure+Archive+Storage)

(linkedin.com/shareArticle?mini=true&url=https%3A%2F%2Fazure.microsoft.com%2Fblog%2Fcloud-storage-now-more-affordable-general-availability-of-azure-archive-storage%2F)

Posted on 13 December, 2017

Kumail Hussain

Principal Program Manager, Azure Storage

As of August 5, 2019, Azure Archive Storage dropped prices by up to 50 percent in some regions. For more information, please refer to the blog post, "We're making Azure Archive Storage better with new lower pricing. (<https://azure.microsoft.com/en-us/blog/we-re-making-azure-archive-storage-better-with-new-lower-pricing/>)"

Today we're excited to announce the general availability of **Archive Blob Storage** starting at an industry leading price of **\$0.002 per gigabyte per month!** Last year, we launched Cool Blob Storage (<https://azure.microsoft.com/blog/introducing-azure-cool-storage/>) to help customers reduce storage costs by tiering their infrequently accessed data to the Cool tier. Organizations can now reduce their storage costs even further by storing their rarely accessed data in the Archive tier. Furthermore, we're also excited to announce the general availability of **Blob-Level Tiering**, which enables customers to optimize storage costs by easily managing the lifecycle of their data across these tiers at the object level.

From startups to large organizations, our customers in every industry have experienced exponential growth of their data. A significant amount of this data is rarely accessed but must be stored for a long period of time to meet either business continuity or compliance requirements; think employee data, medical records, customer information, financial records, backups, etc. Additionally, recent and coming advances in artificial intelligence and data analytics are unlocking value from data that might have previously been discarded. Customers in many industries want to keep more of these data sets for a longer period but need a scalable and cost-effective solution to do so.

"We have been working with the Azure team to preview Archive Blob Storage for our cloud archiving service for several months now. I love how easy it is to change the storage tier on an existing object via a single API. This allows us to build Information Lifecycle Management into our application logic directly and use Archive Blob Storage to significantly decrease our total Azure Storage costs."

-Tom Inglis, Director of Enabling Solutions at BP

(<https://azurecomcdn.azureedge.net/mediahandler/acomblog/media/Default/blog/f552e0dc-0cd8-4955-816f-ebadf7f7df80.jpg>).



Azure Archive Blob Storage

Azure Archive Blob storage is designed to provide organizations with a low cost means of delivering durable, highly available, secure cloud storage for rarely accessed data with flexible latency requirements (on the order of hours). See [Azure Blob Storage: Hot, cool, and archive tiers](https://docs.microsoft.com/azure/storage/blobs/storage-blob-storage-tiers) (<https://docs.microsoft.com/azure/storage/blobs/storage-blob-storage-tiers>) to learn more.

Archive Storage characteristics include:

- **Cost-effectiveness:** Archive access tier is our lowest priced storage offering for long-term storage which is rarely accessed. Preview pricing will continue through February 2018. For new pricing effective March 1, 2018, see [Azure Block Blob Pricing](https://azure.microsoft.com/pricing/details/storage/blobs/) (<https://azure.microsoft.com/pricing/details/storage/blobs/>).
- **Seamless Integration:** Customers use the same familiar operations on objects in the Archive tier as on objects in the Hot and Cool access tiers. This will enable customers to easily integrate the new access tier into their applications.
- **Durability:** All access tiers including Archive are designed to offer the same [high durability](https://docs.microsoft.com/en-us/azure/storage/common/storage-redundancy) (<https://docs.microsoft.com/en-us/azure/storage/common/storage-redundancy>) that customers have come to expect from Azure Storage with the same data replication options available today.
- **Security:** All data in the Archive access tier is automatically encrypted at rest using 256-bit [AES encryption](https://en.wikipedia.org/wiki/Advanced_Encryption_Standard) (https://en.wikipedia.org/wiki/Advanced_Encryption_Standard), one of the strongest block ciphers available.
- **Global Reach:** Archive Storage is available today in 14 regions – North Central US, South Central US, East US, West US, East US 2, Central US, West US 2, West Central US, North Europe, West Europe, Korea

NYSCentral, Korea&South, Central India, and South India. We're continuing to expand the list of available regions, see the latest Archive Storage regional availability (<https://azure.microsoft.com/en-us/global-infrastructure/services/?products=storage>).

Blob-Level Tiering: easily optimize storage costs without moving data

To simplify data lifecycle management, we now allow customers to tier their data at the object level. Customers can easily change the access tier of a single object among the Hot, Cool, or Archive tiers as usage patterns change, without having to move data between accounts. Blobs in all three access tiers can co-exist within the same account.

Flexible management

Archive Storage and Blob-Level Tiering are available on both new and existing Blob Storage and General Purpose v2 (GPv2) accounts. GPv2 accounts are a new account type that support all our latest features, while offering support for Block Blobs, Page Blobs, Files, Queues, and Tables. Customers with General Purpose v1 (GPv1) accounts can easily convert their accounts to a General Purpose v2 account through a simple 1-click step (Blob Storage account conversion support coming soon). GPv2 accounts have a [different pricing model](https://azure.microsoft.com/pricing/details/storage/) (<https://azure.microsoft.com/pricing/details/storage/>) than GPv1 accounts, and customers should review it prior to using GPv2 as it may change their bill. See [Azure Storage Options](https://aka.ms/azure-storage-accounts) (<https://aka.ms/azure-storage-accounts>) to learn more about GPv2, including how and when to use it.

A user may access Archive and Blob-Level Tiering via the Azure portal (Figure 1), PowerShell, and CLI tools and REST APIs, .NET (Figure 2), Java, Python, or Node.js client libraries.

The screenshot shows the Azure portal interface for managing blobs in a container. On the left, the 'examplecontainer' blade displays a list of blobs: 'Confirmation dialog.PNG', 'Multi-Item Selected.PNG', and 'One Item Selected.PNG'. The first blob is selected. On the right, the 'Blob properties' blade is open for 'Confirmation dialog.PNG'. It shows details like Name: Confirmation dialog.PNG, URL: https://xarchivetestwithperms.blob.core.windows.net/confirmationdialog/Confirmation%20dialog.PNG, Last Modified: 8/4/2017 8:35:17 AM, and Type: Block blob. A dropdown menu labeled 'ACCESS TIER' is open, listing 'Cool', 'Hot', 'Cool' (which is highlighted), and 'Archive'. The 'Hot' option is also visible in the dropdown.

(<https://azurecomcdn.azureedge.net/mediahandler/acomblog/media/Default/blog/515a4e17-399a-4580-b6b5-e479e04b75dd.png>)

Figure 1: Set blob access tier in portal

```
CloudBlockBlob blob = (CloudBlockBlob)items;
blob.SetStandardBlobTier(StandardBlobTier.Archive);
```

RECEIVED NYSCEF: 10/14/2022

Figure 2: Set blob access tier using .NET client library

Partner Integration

We integrate with a broad ecosystem of partners to jointly deliver solutions to our customers. The following partners support Archive Storage:

COMMVAULT

(<https://azurecomcdn.azureedge.net/mediahandler/acomblog/media/Default/blog/1952ea0e-158b-4112-b593-73c7497b4091.png>) Commvault's Windows/Azure Centric software solution enables a single solution for storage-agnostic, heterogeneous enterprise data management. Commvault's native support for Azure, including being one of the first Windows/ISV to be "Azure Certified" has been a key benefit for customers considering a Digital Transformation to Azure. Commvault remains committed to continuing our integration and compatibility efforts with Microsoft, befitting a close relationship between the companies that has existed for over 17 years. **This includes quick, cost effective and efficient movement of data to Azure while enabling indexing such that our customers can proactively use the data we send to Azure, including "Azure Archive".** With this new Archive Storage offering, Microsoft again makes significant enhancements to their Azure offering and we expect that this service will be an important driver of new and expanding opportunities for both Commvault and Microsoft.



NetApp (<https://azurecomcdn.azureedge.net/mediahandler/acomblog/media/Default/blog/394440fc-faf4-49b2-9be7-0a8d7c6e8ec2.png>) NetApp® AltaVault™ cloud-integrated storage enables customers to tap into cloud economics and securely backup data to Microsoft Azure cloud storage at up to 90% lower cost compared with on-premises solutions. AltaVault's modern storage architecture optimizes data using class-leading deduplication, compression, and encryption. Optimized data is written to Azure blob storage, reducing WAN bandwidth requirements and ensuring maximum data security. **By adding Day 1 support for Azure Archive storage (<https://blog.netapp.com/lower-your-cloud-storage-costs-with-azure-archive-and-altavault/>), AltaVault provides organizations access to the most cost-effective Azure blob storage tier, significantly driving down costs for rarely accessed long term backup and archive datasets.** Try AltaVault's free 90-day trial and see how easy it is to leverage Microsoft Azure Archive cloud storage today.

HubStor™

(<https://azurecomcdn.azureedge.net/mediahandler/acomblog/media/Default/blog/dfebde6c-d6ed-4c6a-b9b9-e1907e627d83.png>) HubStor is a cloud archiving platform that converges long-term retention and data protection for on-premises file servers, Office 365, email, and other sources of unstructured data content. Delivered as Software-as-a-Service (SaaS) exclusively on the Azure cloud platform, IT teams are adopting HubStor to understand, secure, and manage large volumes of data in Azure with policies for

NYS classification indexing, WORM retention, deletion, and tiering. As detailed in this post NYSCEF: 10/14/2022

(<https://blog.hubstor.net/how-to-use-azure-archive-blob-storage-for-long-term-data-retention>),

customers can now apply HubStor's built-in file analytics and storage tiering policies with the new Azure Archive Blob Storage tier to place the right data on the optimal tier at the best time in the information lifecycle. Enterprise Strategy Group recently completed a lab validation report on HubStor which you can download here (<https://www.hubstor.net/download/15260/>).



(<https://azurecomcdn.azureedge.net/mediahandler/acomblog/media/Default/blog/89272565-1a9c-4391-b1ba-1ab8960cc68c.png>) The purpose of CloudBerry Backup for Microsoft Azure is automating data upload to Microsoft Azure cloud storage. It is able to compress and encrypt the data with a user-defined password prior to the upload. It then securely transfers it to the cloud either on schedule or in real time. CloudBerry Backup also comes with file-system and image-based backup, SQL Server and MS Exchange support, as well as flexible retention policies and incremental backup. **CloudBerry Backup now supports Microsoft Azure Archive Blob Storage** (<https://www.cloudberrylab.com/blog/meet-microsoft-azure-archive-blob-storage/>). **for storing backup and archival data.**



(<https://azurecomcdn.azureedge.net/mediahandler/acomblog/media/Default/blog/b2bd9d5c-8f68-4880-a4ad-cd8397e95e5c.png>) Archive2Azure, the intelligent data management and compliance archiving solution, provides customers a native Azure archiving application. Archive2Azure enables companies to provide automated retention, indexing on demand, encryption, search, review, and production for long term archiving of their compliance, active, low-touch, and inactive data from within their own Azure tenancy. This pairing of the Azure Cloud with Archive2Azure's archiving and data management capabilities provides companies with the cloud-based security and information management they have long sought. With the general availability of Azure's much anticipated Archive Storage offering, the needed security and lower cost to archive and manage data for extended periods is now possible. **With the availability of the new Archive Storage offering, Archive2Azure can now offer Azure's full range of storage tiers** (<https://blog.archive360.com/to-tier-or-not-to-tier-is-no-longer-the-question>). **providing users a wide choice of storage performance and cost.**

COHESITY

(<https://azurecomcdn.azureedge.net/mediahandler/acomblog/media/Default/blog/bc520de0-aef8-4858-9fbf-99df2a46237b.png>) **[Archive support coming soon]** Cohesity delivers the world's first hyper-converged storage system for enterprise data. Cohesity consolidates fragmented, inefficient islands of secondary storage into an infinitely expandable and limitless storage platform that can run both on-premises and in the public cloud. Designed with the latest web-scale distributed systems technology, Cohesity radically simplifies existing backup, file shares, object, and dev/test storage silos by creating a unified, instantly-accessible storage pool. **The Cohesity platform will support Azure Archive Storage** (<https://www.cohesity.com/blog/cohesity-supports-new-microsoft-azure-archive-blob-storage-tier-hybrid-use-cases/>). **for the following customer use cases:** **(i) long-term data retention for infrequently accessed data that require cost effective lowest priced blob storage, (ii) blob-level**

NYS INDEX NO. 654332/2020 FILED 10/14/2022

tiering functionality among Hot, Cool and Archive tiers, and (iii) ease of recovery of data from cloud back to on-premise independent of which Azure blob tier the data is in. Note that Azure Blob storage can be easily registered and assigned via Cohesity's policy-based administration portal to any data protection workload running on the Cohesity platform.



(<https://azurecomcdn.azureedge.net/mediahandler/acomblog/media/Default/blog/834c6b1a-595a-4326-8c3e-cb05e7d846b5.png>) Igneous Systems delivers the industry's first secondary storage system built to handle massive file systems. Offered as-a-Service and built using a cloud-native architecture, Igneous Hybrid Storage Cloud provides a modern, scalable approach to management of unstructured file data across datacenters and public cloud, without the need to manage infrastructure. **Igneous supports backup and long-term archiving of unstructured file data to Azure Archive Blob Storage** (<https://inside.igneous.io/announcing-easy-tiering-and-replication-to-microsoft-azure-archive-blob-storage-with-igneous>), **enabling organizations to replace legacy backup software and targets with a hybrid cloud approach.**

Announcements (/en-au/blog/topics/announcements/) Storage, Backup & Recovery (/en-au/blog/topics/storage-backup-and-recovery/) Management (/en-au/blog/topics/management/)

[Subscribe \(//azurecomcdn.azureedge.net/en-au/blog/feed/\)](//azurecomcdn.azureedge.net/en-au/blog/feed/)

Explore

See where we're heading. Check out forthcoming changes to Azure products

[Azure updates \(/en-au/updates/\)](#)

Let us know if you have any additional questions about Azure

[Ask questions \(https://aka.ms/azureqa\)](https://aka.ms/azureqa)

Topics

[Announcements \(/en-au/blog/topics/announcements/\)](#) (2352)

[Artificial Intelligence \(/en-au/blog/topics/artificial-intelligence/\)](#) (280)

[Azure Maps \(/en-au/blog/topics/azure-maps/\)](#) (34)

[Azure Marketplace \(/en-au/blog/topics/azure-marketplace/\)](#) (154)

[Azure Stream Analytics \(/en-au/blog/topics/azure-stream-analytics/\)](#) (38)

[Big Data \(/en-au/blog/topics/big-data/\)](#) (664)

[Blockchain \(/en-au/blog/topics/blockchain/\)](#) (89)

[Business Intelligence \(/en-au/blog/topics/business-intelligence/\)](#) (122)

[Cloud Strategy \(/en-au/blog/topics/cloud-strategy/\)](#) (816)

[Cognitive Services \(/en-au/blog/topics/cognitive-services/\)](#) (129)

[Customer Enablement \(/en-au/blog/topics/customerenablement/\)](#) (26)

[Data Science \(/en-au/blog/topics/datascience/\)](#) (116)

[Data Warehouse \(/en-au/blog/topics/data-warehouse/\)](#) (223)

[Database \(/en-au/blog/topics/database/\)](#) (644)

[Developer \(/en-au/blog/topics/developer/\)](#) (1219)

[DevOps \(/en-au/blog/topics/devops/\)](#) (98)

[Events \(/en-au/blog/topics/events/\)](#) (275)

[Government \(/en-au/blog/topics/government/\)](#) (79)

[Hybrid \(/en-au/blog/topics/hybrid/\)](#) (91)

[Identity & Access Management \(/en-au/blog/topics/identity-access-management/\)](#) (90)

[Internet of Things \(/en-au/blog/topics/internet-of-things/\)](#) (397)

[IT Pro \(/en-au/blog/topics/it-pro/\)](#) (603)

[Last week in Azure \(/en-au/blog/topics/last-week-in-azure/\)](#) (93)

[Machine Learning \(/en-au/blog/topics/machine-learning/\)](#) (69)

[Management \(/en-au/blog/topics/management/\)](#) (414)

[Media Services & CDN \(/en-au/blog/topics/media-services/\)](#) (207)

[Migration \(/en-au/blog/topics/migration/\)](#) (54)

[Mobile \(/en-au/blog/topics/mobile/\)](#) (160)

[Monitoring \(/en-au/blog/topics/monitor/\)](#) (158)

[Networking \(/en-au/blog/topics/networking/\)](#) (274)

[Partner \(/en-au/blog/topics/partner/\)](#) (171)

[Serverless](#) (/en-au/blog/topics/serverless/) (81)

[Storage, Backup & Recovery](#) (/en-au/blog/topics/storage-backup-and-recovery/) (712)

[Supportability](#) (/en-au/blog/topics/supportability/) (47)

[Updates](#) (/en-au/blog/topics/updates/) (593)

[Virtual Machines](#) (/en-au/blog/topics/virtual-machines/) (766)

[Web](#) (/en-au/blog/topics/web/) (379)

Articles by date

[September 2022](#) (/en-au/blog/2022/09/).

[August 2022](#) (/en-au/blog/2022/08/).

[July 2022](#) (/en-au/blog/2022/07/).

[June 2022](#) (/en-au/blog/2022/06/).

[May 2022](#) (/en-au/blog/2022/05/).

[April 2022](#) (/en-au/blog/2022/04/).

[Full archive](#) (/en-au/blog/archives/)